

EXPERIMENT NO: 3

AIM : To understand the Kubernetes Cluster Architecture, install and Spin Up a Kubernetes Cluster on Linux Machines/Cloud Platforms.

STEPS :**1. Create Instances:**

We initiated the creation of three virtual machines or instances, naming them Master, Worker1, and Worker2. These instances will act as the nodes in our Kubernetes cluster.

Master	i-0767a02f53056b254	Running	t3.small	Initializing	View alarms +
worker-1	i-0a98404682c2bf690	Running	t3.small	Initializing	View alarms +
worker-2	i-0a9ea3e263873d151	Running	t3.small	Initializing	View alarms +

2. Install Docker:

On each instance, execute the following commands to gain superuser privileges and install Docker:

```
sudo su
```

```
yum install docker -y
```

This step ensures that Docker, a container runtime, is installed on Master, Worker1, and Worker2. Docker is essential for running the containerized applications that Kubernetes will orchestrate.

```
[ec2-user@ip-172-31-93-226 ~]$ sudo su
[root@ip-172-31-93-226 ec2-user]# yum install docker -y
Last metadata expiration check: 0:07:12 ago on Fri Sep 13 11:58:42 2024.
Dependencies resolved.
=====
Package Architecture Version Repo
sitortory Size
=====
Installing:
docker x86_64 25.0.6-1.amzn2023.0.2 amaz
onlinux 44 M
Installing dependencies:
containerd x86_64 1.7.20-1.amzn2023.0.1 amaz
onlinux 35 M
iptables-libs x86_64 1.8.8-3.amzn2023.0.2 amaz
onlinux 401 k
iptables-nft x86_64 1.8.8-3.amzn2023.0.2 amaz
onlinux 183 k
```

3.Start Docker:

Start the Docker service on each instance to ensure it is running:

```
systemctl start docker
```

- This command activates the Docker service on Master, Worker1, and Worker2, enabling them to run and manage containers.

```
[root@ip-172-31-93-226 ec2-user]# systemctl start docker  
[root@ip-172-31-93-226 ec2-user]#
```

4.Install kubeadm:

Prepare the system for Kubernetes installation by configuring SELinux and adding the Kubernetes repository:

```
sudo setenforce 0  
sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config  
cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo  
[kubernetes]  
name=Kubernetes  
baseurl=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/  
enabled=1  
gpgcheck=1  
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/repodata/repomd.xml.key  
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni  
EOF
```

Install Kubernetes components:

```
sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes  
sudo systemctl enable --now kubelet
```

These steps configure SELinux to run in permissive mode, add the Kubernetes repository, and install Kubernetes tools (kubelet, kubeadm, and kubectl) on Master, Worker1, and Worker2. The kubelet service is also enabled and started to ensure it runs on boot.

```
Installed:
  conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64      cri-tools-1.31.1-150500.1.1.x86_64      kubeadm-1.31.1-150500.1.1.x86_64
  kubect1-1.31.1-150500.1.1.x86_64                kubelet-1.31.1-150500.1.1.x86_64        kubernetes-cni-1.5.1-150500.1.1.x86_64
  libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64  libnetfilter_cttimeout-1.0.0-19.amzn2023.0.2.x86_64  libnetfilter_queue-1.0.5-2.amzn2023.0.2.x86_64

Complete!
[root@ip-172-31-84-46 ec2-user]# sudo systemctl enable --now kubelet
Created symlink /etc/systemd/system/multi-user.target.wants/kubelet.service → /usr/lib/systemd/system/kubelet.service.
```

5. Confirm Repository:

Verify that the Kubernetes repository is correctly configured by listing all enabled repositories:

bash

Copy code

yum repo list

Execute this command on Master, Worker1, and Worker2 to confirm that the Kubernetes repository has been added and is available for package installations.

```
[root@ip-172-31-84-46 ec2-user]# yum repolist
repo id                                repo name
amazonlinux                            Amazon Linux 2023 repository
kernel-livepatch                       Amazon Linux 2023 Kernel Livepatch repository
kubernetes                             Kubernetes
[root@ip-172-31-84-46 ec2-user]#
```

6.Initialize kubeadm on Master Node:

On the Master node, initialize the Kubernetes cluster using kubeadm. This process sets up the Kubernetes control plane and generates commands for joining worker nodes:

Copy the commands displayed in the output of the initialization process to configure permissions and obtain the join token. This includes a join command link needed for worker nodes to connect to the master.

```
Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config

Alternatively, if you are the root user, you can run:

export KUBECONFIG=/etc/kubernetes/admin.conf

You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
https://kubernetes.io/docs/concepts/cluster-administration/addons/

Then you can join any number of worker nodes by running the following on each as root:

kubeadm join 172.31.84.46:6443 --token jl06ac.t7cdzxf0x5eddmsl \
--discovery-token-ca-cert-hash sha256:4a152b913ee4b60dc2126d55f631b86d0dafb7d58132416c4f32f0668ac553be
[root@ip-172-31-84-46 ec2-user]#
```

In the screenshot you can see the commands written in the 3rd 4th and 5th line

Copy that command , this command is used to add right permission to the user

Also copy the 7th line , here it is the credential for the user .

Also copy the last 2 lines it is a link used to join the nodes

7.Join Worker Nodes:

On Worker1 and Worker2, use the join command provided by the Master node to connect them to the Kubernetes cluster:

```
kubeadm join 172.31.84.46:6443 --token jl06ac.t7cdzxf0x5eddmsl \
--discovery-token-ca-cert-hash
sha256:4a152b913ee4b60dc2126d55f631b86d0dafb7d58132416c4f32f0668ac553be
```

This command allows Worker1 and Worker2 to register themselves with the Master node and become part of the Kubernetes cluster.

Verify Nodes:

```
kubectl get nodes
```

```
[root@ip-172-31-84-46 ec2-user]# kubectl get node
NAME                                STATUS    ROLES    AGE   VERSION
ip-172-31-84-46.ec2.internal       NotReady control-plane 56s   v1.31.1
[root@ip-172-31-84-46 ec2-user]# kubectl get node
The connection to the server 172.31.84.46:6443 was refused - did you specify the right host or port?
[root@ip-172-31-84-46 ec2-user]# kubectl get node
NAME                                STATUS    ROLES    AGE   VERSION
ip-172-31-84-46.ec2.internal       NotReady control-plane 5m1s   v1.31.1
[root@ip-172-31-84-46 ec2-user]#
```

```

#
##### Amazon Linux 2023
##### \
##### \
##### |
##### /
V~' -> https://aws.amazon.com/linux/amazon-linux-2023

~/m'

Last login: Fri Sep 13 12:04:33 2024 from 18.206.107.28
[ec2-user@ip-172-31-87-149 ~]$ sudo su
[root@ip-172-31-87-149 ec2-user]# kubeadm join 172.31.84.46:6443 --token qq448c.8c0zquwywz3egt0b --discovery-token-ca-cert-hash sha256:49cc770e646aab704883a3d9fb30e6146fdd83e782ebd5ec3194dacee59f2257
[preflight] Running pre-flight checks
[WARNING FileExisting-socat]: socat not found in system path
[WARNING FileExisting-tc]: tc not found in system path
```

The Cloud shell didn't proceed further.

We created and configured three instances (Master, Worker1, and Worker2), installed Docker and Kubernetes on all nodes, and initialized the Kubernetes control plane on the Master node. Despite providing the join command to Worker1 and Worker2, they failed to connect to the Master node.

No Running Containers: The `docker ps -a` command shows no containers, which means the Kubernetes components are not running in Docker.

CrashLoopBackOff: There are errors indicating that the containers for Kubernetes components are restarting repeatedly but failing to start properly.

Network Plugin Not Ready: The error message "Network plugin returns error: cni plugin not initialized" suggests issues with the network plugin required for Kubernetes.